



Securing Kubernetes

Protecting Your Cloud: A Simple Guide to Securing Kubernetes

Outline

01 The Problem

Current Kubernetes security challenges and attack vectors

02 Defense Strategy

Core security strategies for protecting a Kubernetes cluster across multiple levels

03 Runtime Defense

Tools and practices for detecting and mitigating threats in real-time.

04 Next Steps

Call to action for continuous security improvement.

01

The Problem

What is Kubernetes?

Before Kubernetes

- Scaling is PAINFUL
- Dev/Test/Prod Environment inconsistent
- Downtime is COMMON

After Kubernetes

- Scaling is AUTOMATED
- Consistent environment
- Downtime is minimized



What is Kubernetes?

Kubernetes is an open source container orchestration engine for automating deployment, scaling, and management of containerized applications

> Terms

Master Node: Machine that run control plane

Control Plane: Services that manage cluster

Node: Worker Machine running pods, managed by control plane

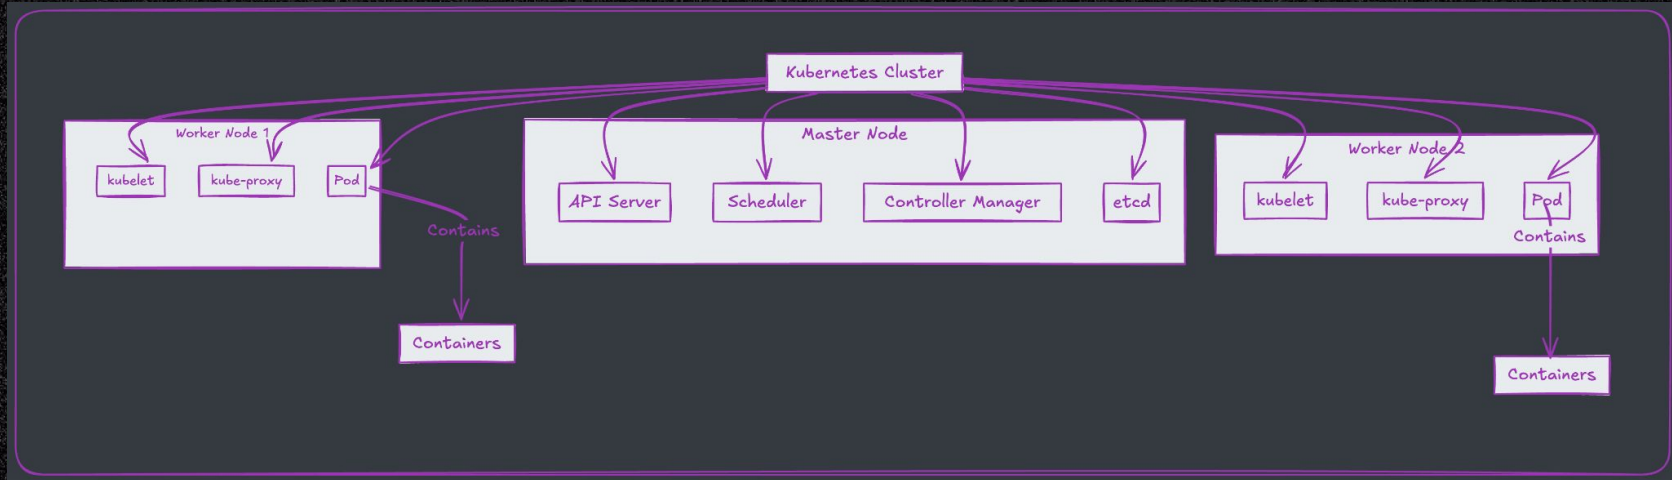
Kubelet: Agent that communicates with the control plane to receive command from control plane and report node states

Kube-proxy: Runs on each node, managing service communication by forwarding traffic to the app backends and supporting cluster IPs via environment variables or optional DNS.

Pod: Smallest deployable unit, contains one or more containers



Kubernetes Cluster



Autoscaling



The Cost of Breach

2021 – \$1.8M compute costs

Cryptojacking campaign hijacked 2,400+ public clusters through weak kubelet authentication (anonymous auth enabled)

2021 – 500K AWS Credentials Stolen

Malicious Helm chart "redis-optimized" stole AWS credentials.

2024

CVE-2024-9486 – Kubernetes Image Builder versions \leq v0.137 allows attackers to gain root access to nodes using VM images created with the Proxmox provider due to enabled default credentials during the image build process

Source:

- <https://www.microsoft.com/en-us/security/business/security-intelligence-report>
- <https://www.tenable.com/cve/CVE-2024-9486>

The Threat Landscape:

2025's Attack Playbook




63% from misconfigurations

37% (supply chain attack, vulnerability on images, insider attack, etc)

02

Defense Strategy

Defense Strategy



CONTROL PLANE

The control plane is the brain of the Kubernetes cluster

- **Restrict Access:** Use bastion hosts or VPNs to limit access to nodes hosting control plane components, ensuring only authorized personnel can interact with them.
- **Enforce TLS Encryption:** Ensure all communication between control plane components and API servers is encrypted using Transport Layer Security (TLS) to prevent man-in-the-middle attacks
- **Role-Based Access Control (RBAC):** Implement RBAC policies to enforce least privilege access. Regularly audit permissions and test configurations by impersonating users or user groups via user impersonation

Tools: kube-bench, cert-manager, hashicorp vault

kube-bench

kube-bench is a tool that checks whether Kubernetes is deployed securely by running the checks documented in the CIS Kubernetes Benchmark.

```
$ kubectl apply -f job.yaml
job.batch/kube-bench created

$ kubectl get pods
NAME                                READY   STATUS             RESTARTS   AGE
kube-bench-j76s9    0/1     ContainerCreating   0           3s

# Wait for a few seconds for the job to complete
$ kubectl get pods
NAME                                READY   STATUS             RESTARTS   AGE
kube-bench-j76s9    0/1     Completed           0          11s

# The results are held in the pod's logs
kubectl logs kube-bench-j76s9
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 API Server
...
```


kube-bench

```
[INFO] 4 Worker Node Security Configuration
[INFO] 4.1 Worker Node Configuration Files
[PASS] 4.1.1 Ensure that the kubelet service file permissions are set to 644 or more restrictive (Automated)
[PASS] 4.1.2 Ensure that the kubelet service file ownership is set to root:root (Automated)
[PASS] 4.1.3 If proxy kubeconfig file exists ensure permissions are set to 644 or more restrictive (Manual)
[PASS] 4.1.4 Ensure that the proxy kubeconfig file ownership is set to root:root (Manual)
[PASS] 4.1.5 Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive (Automated)
[PASS] 4.1.6 Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root (Manual)
[PASS] 4.1.7 Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Manual)
[PASS] 4.1.8 Ensure that the client certificate authorities file ownership is set to root:root (Manual)
[PASS] 4.1.9 Ensure that the kubelet --config configuration file has permissions set to 644 or more restrictive (Automated)
[PASS] 4.1.10 Ensure that the kubelet --config configuration file ownership is set to root:root (Automated)
[INFO] 4.2 Kubelet
[PASS] 4.2.1 Ensure that the anonymous-auth argument is set to false (Automated)
[PASS] 4.2.2 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)
[PASS] 4.2.3 Ensure that the --client-ca-file argument is set as appropriate (Automated)
[PASS] 4.2.4 Ensure that the --read-only-port argument is set to 0 (Manual)
[PASS] 4.2.5 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Manual)
[FAIL] 4.2.6 Ensure that the --protect-kernel-defaults argument is set to true (Automated)
[PASS] 4.2.7 Ensure that the --make-iptables-util-chains argument is set to true (Automated)
[PASS] 4.2.8 Ensure that the --hostname-override argument is not set (Manual)
[WARN] 4.2.9 Ensure that the --event-qps argument is set to 0 or a level which ensures appropriate event capture (Manual)
[WARN] 4.2.10 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Manual)
[PASS] 4.2.11 Ensure that the --rotate-certificates argument is not set to false (Manual)
[PASS] 4.2.12 Verify that the RotateKubeletServerCertificate argument is set to true (Manual)
[WARN] 4.2.13 Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Manual)
```


Node

Nodes are the worker machines that run your applications

- **Minimize Attack Surface:** Disable unnecessary services and ports on nodes to reduce exposure to potential threats
- **Regular Updates:** Keep node operating systems and container runtimes up-to-date with the latest patches to mitigate vulnerabilities
- **Immutable Infrastructure:** Use immutable infrastructure practices where possible, ensuring nodes are replaced rather than patched in-place to maintain consistency and reduce drift
- **Node Isolation:** Use dedicated nodes for sensitive workloads and apply taints/tolerations to ensure proper workload placement

Tools: kube-hunter

Network

Network security ensures secure communication within and outside the cluster:

- **Network Policies:** Enforce Kubernetes-native network policies to restrict traffic between pods and namespaces, preventing lateral movement in case of compromise
- **Encrypt Traffic:** Use mutual TLS (mTLS) for service-to-service communication to ensure data confidentiality and integrity
- **Ingress/EGress Controls:** Configure ingress controllers and egress gateways to monitor and filter incoming/outgoing traffic, blocking unauthorized connections

Tools: Calico

Calico

Manage network policy

It runs as daemonset, ensuring
each node run calico pod

\$ kubectl create -f
<https://raw.githubusercontent.com/projectcalico/calico/v3.29.2/manifests/custom-resources.yaml>

Use domain names in a global network set

In this method, you create a **GlobalNetworkSet** with the allowed destination domain names in the `allowedEgressDomains` field. Then, you create a **GlobalNetworkPolicy** with a `destination.selector` that matches that `GlobalNetworkSet`.

In the following example, the allowed egress domains (`api.alice.com` and `*.example.com`) are specified in the `GlobalNetworkSet`.

```
apiVersion: projectcalico.org/v3
kind: GlobalNetworkSet
metadata:
  name: allowed-domains-1
  labels:
    color: red
spec:
  allowedEgressDomains:
    - api.alice.com
    - '*.example.com'
```


Pods

Pods are the smallest deployable units in Kubernetes

- **Least Privilege Principle:** Run containers as non-root users and avoid privileged mode unless absolutely necessary
- **Image Security:** Use trusted base images from verified sources and scan them regularly for vulnerabilities before deployment
- **Resource Limits:** Set CPU/memory limits to prevent resource exhaustion attacks, such as denial-of-service (DoS)
- **Security Contexts:** Define pod security contexts to enforce security settings like read-only filesystems, seccomp profiles, and AppArmor/SELinux policies

Tools: Trivy, checkov, kyverno

Kyverno

Allows you to manage and enforce policies on your Kubernetes clusters

bash

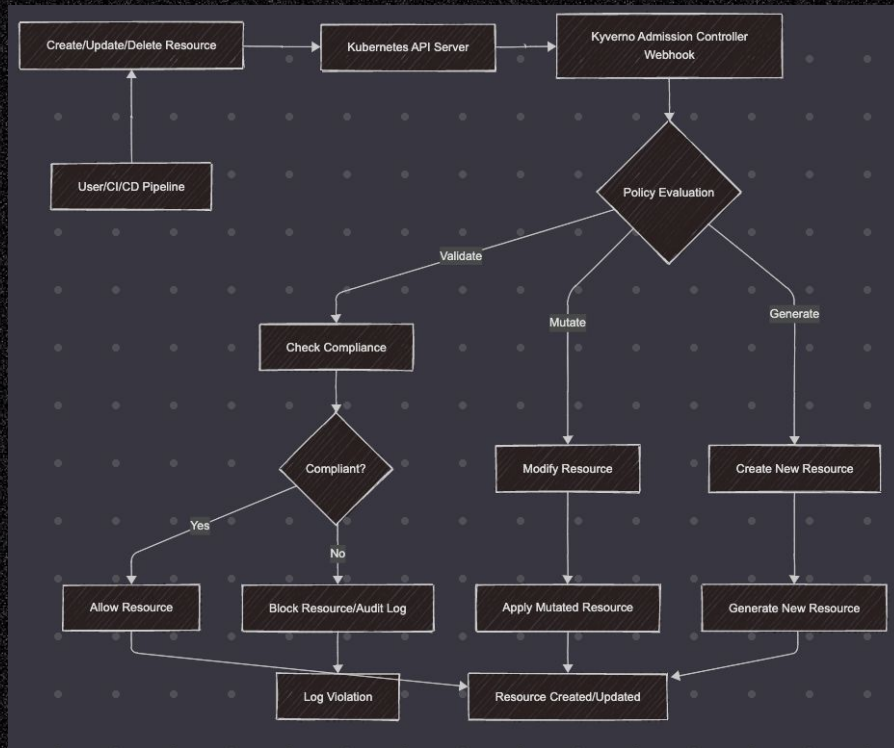
```
1 helm repo add kyverno https://kyverno.github.io/kyverno/  
2 helm repo update  
3 helm install kyverno kyverno/kyverno --namespace kyverno --create-namespace
```

Copy

Example non-compliance pod

Error from server: admission webhook

"validate.kyverno.svc" denied the request: All Pods must have the 'app' label.



kube-hunter

Actively tests your cluster for vulnerabilities, simulating attacks.

```
bash
```

```
1 kube-hunter --remote <your-cluster-ip-or-domain>
```

Active Hunters:

* Foothold Via Secure Kubelet Port

Attempts to demonstrate that a malicious actor can establish foothold into the cluster via a container abusing the configuration of the kubelet's secure port: authentication-auth=false.

* Malicious Intent Via Secure Kubelet Port

Attempts to demonstrate that a malicious actor can leverage existing privileged containers exposed via the kubelet's secure port, due to anonymous auth enabled misconfiguration, such that a process can be started or modified on the host.

* Kubelet Run Hunter

Executes uname inside of a random container

* Kubelet Container Logs Hunter

Retrieves logs from a random container

* Kubelet System Logs Hunter

Retrieves commands from host's system audit

* Azure SPN Hunter

Gets the azure subscription file on the host by executing inside a container

* API server hunter

Accessing the api server might grant an attacker full control over the cluster

* Arp Spoof Hunter

Checks for the possibility of running an ARP spoof attack from within a pod (results are based on the running node)

* DNS Spoof Hunter

Checks for the possibility for a malicious pod to compromise DNS requests of the cluster (results are based on the running node)

* Etcd Remote Access

Checks for remote write access to etcd, will attempt to add a new key to the etcd DB

* Prove /var/log Mount Hunter

Tries to read /etc/shadow on the host by running commands inside a pod with host mount to /var/log

* Build Date Hunter

Hunts when proxy is exposed, extracts the build date of kubernetes

* K8s Version Hunter

Hunts Proxy when exposed, extracts the version



Detect vulnerabilities, misconfigurations, and other security issues in container images.

```
2. bash
bash-3.2$ trivy knqyf263/test-image:1.2.3
2019-05-13T15:19:03.012+0900 INFO Updating vulnerability database...
2019-05-13T15:19:05.983+0900 INFO Detecting Alpine vulnerabilities...
2019-05-13T15:19:05.987+0900 INFO Updating npm Security DB...
2019-05-13T15:19:07.048+0900 INFO Detecting npm vulnerabilities...
2019-05-13T15:19:07.048+0900 INFO Updating pipenv Security DB...
2019-05-13T15:19:08.507+0900 INFO Detecting pipenv vulnerabilities...
2019-05-13T15:19:08.509+0900 INFO Updating bundler Security DB...
2019-05-13T15:19:09.574+0900 INFO Detecting bundler vulnerabilities...
2019-05-13T15:19:09.575+0900 INFO Updating cargo Security DB...
2019-05-13T15:19:10.441+0900 INFO Detecting cargo vulnerabilities...
2019-05-13T15:19:10.441+0900 INFO Updating composer Security DB...
2019-05-13T15:19:11.649+0900 INFO Detecting composer vulnerabilities...

knqyf263/test-image:1.2.3 (alpine 3.7.1)
=====
Total: 26 (UNKNOWN: 0, LOW: 3, MEDIUM: 16, HIGH: 5, CRITICAL: 2)

+-----+-----+-----+-----+-----+-----+
| LIBRARY | VULNERABILITY ID | SEVERITY | INSTALLED VERSION | FIXED VERSION | TITLE |
+-----+-----+-----+-----+-----+-----+
| curl | CVE-2018-14618 | CRITICAL | 7.61.0-r0 | 7.61.1-r0 | curl: NTLM password overflow |
| | | | | | via integer overflow |
| | CVE-2018-16839 | HIGH | | 7.61.1-r1 | curl: Integer overflow leading |
| | | | | | to heap-based buffer overflow in |
| | | | | | Curl_sasl_create_plain_message() |
| | CVE-2019-3822 | | | 7.61.1-r2 | curl: NTLMv2 type-3 header |
| | | | | | stack buffer overflow |
| | CVE-2018-16840 | | | 7.61.1-r1 | curl: Use-after-free when |
| | | | | | closing "easy" handle in |
| | | | | | Curl_close() |
| | CVE-2018-16890 | MEDIUM | | 7.61.1-r2 | curl: NTLM type-2 heap |
| | | | | | out-of-bounds buffer read |
| | CVE-2019-3823 | | | | curl: SMTP end-of-response |
+-----+-----+-----+-----+-----+-----+
```

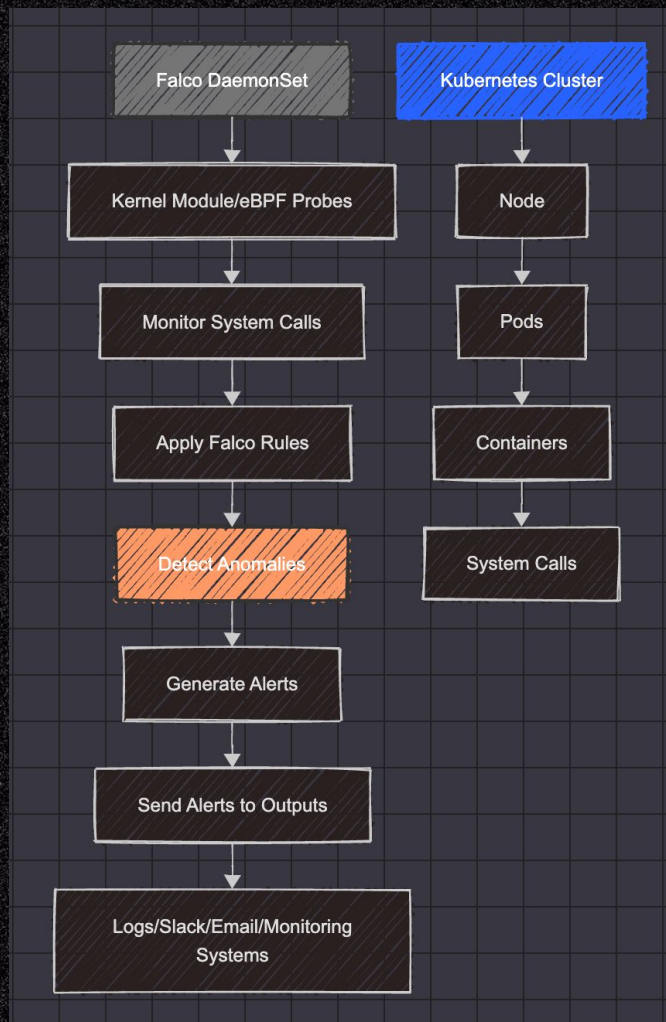

03

Runtime Defense

FALCO

Runtime Security monitoring & detection

```
$ helm repo add falcosecurity  
https://falcosecurity.github.io/charts  
$ helm repo update  
$ helm install --replace falco --namespace falco  
--create-namespace --set tty=true falcosecurity/falco
```



04

Next Steps

Next Steps

High Impact Low Effort

- Run kube-bench to audit (policy misconfiguration, restricted access, etc)
- Enforce Network Policies with Calico
- Regular Update Node OS

Next - Next Steps

- Micro segmentation
- Image scanner
- Runtime defense
- etc

Thank you